

KNN

```
# Import required libraries
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.datasets import fetch_openml
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.preprocessing import StandardScaler
from sklearn.neighbors import KNeighborsClassifier
from sklearn.decomposition import PCA
from sklearn.metrics import accuracy_score, confusion_matrix,
classification_report

# Load the MNIST dataset
print("Loading MNIST dataset...")
mnist = fetch_openml('mnist_784', version=1)
X, y = mnist.data, mnist.target.astype(np.uint8) # Convert labels to
integers

# Split into training (80%) and testing (20%)
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

print(f"Train shape: {X_train.shape}, Test shape: {X_test.shape}")

# Normalize the data (important for kNN)
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)

# Hyperparameter tuning: Find best k value using cross-validation
print("\nFinding best k using cross-validation...")
k_values = [1, 3, 5, 7, 9]
cv_scores = [cross_val_score(KNeighborsClassifier(n_neighbors=k,
n_jobs=-1), X_train, y_train, cv=3).mean() for k in k_values]

best_k = k_values[np.argmax(cv_scores)]
print(f"Best k found: {best_k}")

# Train kNN model with best k
print("\nTraining kNN model...")
knn = KNeighborsClassifier(n_neighbors=best_k, n_jobs=-1)
knn.fit(X_train, y_train)

# Predict on test data
y_pred = knn.predict(X_test)
```

```

# Evaluate model performance
accuracy = accuracy_score(y_test, y_pred)
print(f"\nTest Accuracy: {accuracy:.4f}")

# Confusion matrix
conf_matrix = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(8,6))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap="Blues",
xticklabels=range(10), yticklabels=range(10))
plt.xlabel("Predicted Label")
plt.ylabel("True Label")
plt.title("Confusion Matrix")
plt.show()

# Classification report
print("\nClassification Report:")
print(classification_report(y_test, y_pred))

# Dimensionality Reduction using PCA
print("\nApplying PCA to optimize kNN...")
pca = PCA(n_components=50) # Reduce to 50 features
X_train_pca = pca.fit_transform(X_train)
X_test_pca = pca.transform(X_test)

# Train kNN with PCA-transformed data
knn_pca = KNeighborsClassifier(n_neighbors=best_k, n_jobs=-1)
knn_pca.fit(X_train_pca, y_train)
y_pred_pca = knn_pca.predict(X_test_pca)

# Evaluate model with PCA
pca_accuracy = accuracy_score(y_test, y_pred_pca)
print(f"Test Accuracy with PCA: {pca_accuracy:.4f}")

# Plot some sample predictions
fig, axes = plt.subplots(2, 5, figsize=(10, 5))
axes = axes.ravel()
for i in range(10):
    img = X_test[i].reshape(28, 28) # Reshape to 28x28 image
    axes[i].imshow(img, cmap='gray')
    axes[i].set_title(f"Pred: {y_pred[i]}, Actual: {y_test.iloc[i]}")
    axes[i].axis('off')

plt.tight_layout()
plt.show()


Loading MNIST dataset...
Train shape: (56000, 784), Test shape: (14000, 784)

Finding best k using cross-validation...
Best k found: 3

```

Training kNN model...

Test Accuracy: 0.9465



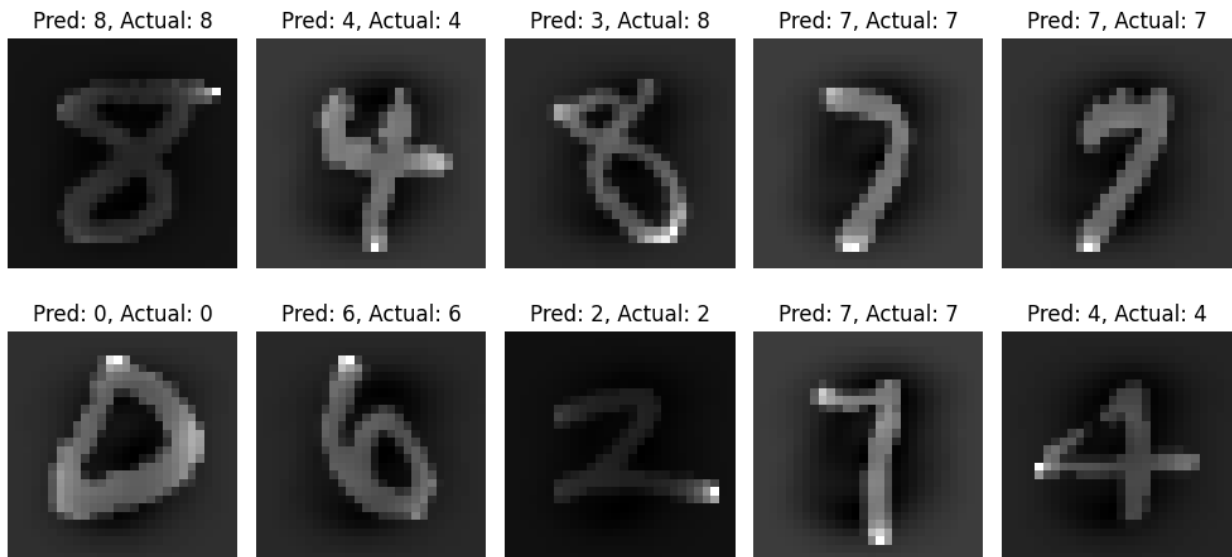
Classification Report:

| | precision | recall | f1-score | support |
|---|-----------|--------|----------|---------|
| 0 | 0.96 | 0.98 | 0.97 | 1343 |
| 1 | 0.95 | 0.99 | 0.97 | 1600 |
| 2 | 0.95 | 0.94 | 0.94 | 1380 |
| 3 | 0.93 | 0.95 | 0.94 | 1433 |
| 4 | 0.95 | 0.94 | 0.94 | 1295 |
| 5 | 0.95 | 0.94 | 0.94 | 1273 |
| 6 | 0.97 | 0.97 | 0.97 | 1396 |
| 7 | 0.94 | 0.93 | 0.93 | 1503 |
| 8 | 0.97 | 0.90 | 0.93 | 1357 |

| | | | | | |
|--------------|---|------|------|------|-------|
| | 9 | 0.90 | 0.92 | 0.91 | 1420 |
| accuracy | | | | 0.95 | 14000 |
| macro avg | | 0.95 | 0.95 | 0.95 | 14000 |
| weighted avg | | 0.95 | 0.95 | 0.95 | 14000 |

Applying PCA to optimize kNN...

Test Accuracy with PCA: 0.9586



CNN

```
# Import required libraries
import numpy as np
import matplotlib.pyplot as plt
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten,
Dense, Dropout, BatchNormalization
from tensorflow.keras.datasets import mnist
from tensorflow.keras.utils import to_categorical

# Load the MNIST dataset
print("Loading MNIST dataset...")
(X_train, y_train), (X_test, y_test) = mnist.load_data()

# Reshape to (28,28,1) since CNN expects 3D input
X_train = X_train.reshape(-1, 28, 28, 1).astype('float32') / 255.0
X_test = X_test.reshape(-1, 28, 28, 1).astype('float32') / 255.0
```

```

# One-hot encode labels
y_train = to_categorical(y_train, 10)
y_test = to_categorical(y_test, 10)

print(f"Train shape: {X_train.shape}, Test shape: {X_test.shape}")

# Build CNN model
model = Sequential([
    Conv2D(32, (3,3), activation='relu', padding='same',
input_shape=(28,28,1)),
    BatchNormalization(),
    MaxPooling2D((2,2)),

    Conv2D(64, (3,3), activation='relu', padding='same'),
    BatchNormalization(),
    MaxPooling2D((2,2)),

    Conv2D(128, (3,3), activation='relu', padding='same'),
    BatchNormalization(),
    MaxPooling2D((2,2)),

    Flatten(),
    Dense(128, activation='relu'),
    Dropout(0.5), # Prevent overfitting
    Dense(10, activation='softmax') # Output layer for 10 classes
])

# Compile the model
model.compile(optimizer='adam', loss='categorical_crossentropy',
metrics=['accuracy'])

# Train the CNN model
print("\nTraining CNN model...")
history = model.fit(X_train, y_train, epochs=10, batch_size=64,
validation_data=(X_test, y_test), verbose=1)

# Evaluate on test set
test_loss, test_acc = model.evaluate(X_test, y_test)
print(f"\nTest Accuracy: {test_acc:.4f}")

# Plot training loss and accuracy
plt.figure(figsize=(12,5))
plt.subplot(1,2,1)
plt.plot(history.history['loss'], label='Train Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.title('Loss Curve')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()

```

```
plt.subplot(1,2,2)
plt.plot(history.history['accuracy'], label='Train Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.title('Accuracy Curve')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()
plt.show()
```

Predict some test images

```
y_pred = np.argmax(model.predict(X_test[:10]), axis=1)
y_actual = np.argmax(y_test[:10], axis=1)
```

Display sample predictions

```
fig, axes = plt.subplots(1, 10, figsize=(15,3))
for i in range(10):
    axes[i].imshow(X_test[i].reshape(28,28), cmap='gray')
    axes[i].set_title(f"Pred: {y_pred[i]}\nActual: {y_actual[i]}")
    axes[i].axis('off')
```

```
plt.show()
```

Loading MNIST dataset...

Train shape: (60000, 28, 28, 1), Test shape: (10000, 28, 28, 1)

Training CNN model...

/usr/local/lib/python3.11/dist-packages/keras/src/layers/convolutional/base_conv.py:107: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.

```
super().__init__(activity_regularizer=activity_regularizer,
**kwargs)
```

Epoch 1/10

938/938 ————— 20s 10ms/step - accuracy: 0.8660 - loss: 0.4513 - val_accuracy: 0.9837 - val_loss: 0.0547

Epoch 2/10

938/938 ————— 4s 4ms/step - accuracy: 0.9774 - loss: 0.0797 - val_accuracy: 0.9894 - val_loss: 0.0331

Epoch 3/10

938/938 ————— 6s 5ms/step - accuracy: 0.9839 - loss: 0.0559 - val_accuracy: 0.9907 - val_loss: 0.0341

Epoch 4/10


938/938 ————— 4s 5ms/step - accuracy: 0.9875 - loss: 0.0422 - val_accuracy: 0.9918 - val_loss: 0.0292

Epoch 5/10

938/938 ————— 4s 4ms/step - accuracy: 0.9889 - loss:

0.0376 - val_accuracy: 0.9908 - val_loss: 0.0322
 Epoch 6/10
 938/938 ————— 5s 5ms/step - accuracy: 0.9896 - loss: 0.0361 - val_accuracy: 0.9922 - val_loss: 0.0297
 Epoch 7/10
 938/938 ————— 5s 4ms/step - accuracy: 0.9910 - loss: 0.0322 - val_accuracy: 0.9899 - val_loss: 0.0341
 Epoch 8/10
 938/938 ————— 4s 4ms/step - accuracy: 0.9926 - loss: 0.0242 - val_accuracy: 0.9928 - val_loss: 0.0327
 Epoch 9/10
 938/938 ————— 5s 4ms/step - accuracy: 0.9929 - loss: 0.0251 - val_accuracy: 0.9915 - val_loss: 0.0399
 Epoch 10/10
 938/938 ————— 4s 4ms/step - accuracy: 0.9936 - loss: 0.0199 - val_accuracy: 0.9914 - val_loss: 0.0343
 313/313 ————— 1s 2ms/step - accuracy: 0.9894 - loss: 0.0392

Test Accuracy: 0.9914



1/1 ————— 1s 627ms/step

